



**DeVeLHOPE**

#codeforimpact

# Introduction to SQL



## SQL and databases

The acronym **SQL** stands for **Structured Query Language**.

A **query** is a request for specific information from a database.

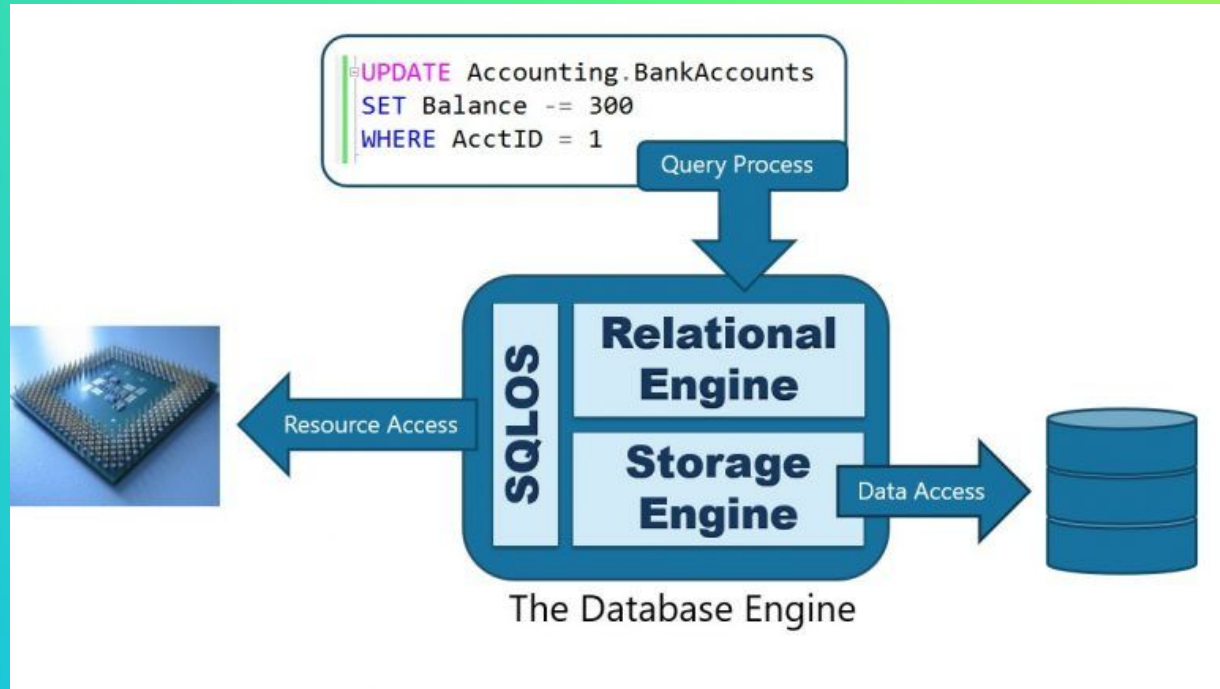
A **database (DB)** is an organised collection of data: for example, SQL-based databases organize data into tables, rows, columns, and indexes in order to facilitate the information management and retrieval.

**SQL** is a language that is used for storing, manipulating and retrieving data in databases.

There are a lot of **different types of databases**, but some of them implement SQL as language.

So, the knowledge of SQL is very important for executing **CRUD (Create, Read, Update, Delete)** operations on databases.

# How a Database works



# Types of databases





## Types of Databases

By data location:

- **Centralized database:** information store in 1 location, typical in big organisations like universities
- **Distributed database:** information is stored across multiple locations
- **Cloud database:** a virtual environment that executes over the cloud computing platform
- **End-user database:** primarily used by a single person, e.g. a spreadsheet on your local computer

By data structure:

- **Graph database:** the connections between data is stored and it's important like the original data
- **NoSQL database:** the hierarchical structure is similar to a file folder system and the stored data is non-relational (that means it isn't structured). NoSQL DBs are suitable for a better scalability
- **Object-oriented database:** data is represented and stored as objects, similarly to OOP principles
- **Relational database:** a **RDB** (*Relational Database*) is based on the *relational data model*, which stores data in the form of rows and columns (that together form a table). RDBs focus on the integrity of data.



## Relational Databases (RDBs)

Considering that the main goal is to learn SQL, we will focus mostly on Relational Databases (e.g. *MySQL*, *PostgreSQL*) because SQL is the language *behind them*.

The **relational data model** is based on the following concepts:

- **Data Domain:** collection of values that a data element can contain
- **Attributes:** they are the basic storage units (also called **columns** or **fields**)
- **Records:** they contain fields that are related (also called **rows** or **tuple**)
- **Relation:** a relation is basically a table structure. Each table has 2 properties:
  - Rows (the records)
  - Columns (the attributes)
- **Relation Key:** each record (row) has an attribute or a combination of attributes that can identify the tuple in a unique way



## Relational Data Model - an example

**Table also called Relation**

Primary Key      Domain  
Ex: NOT NULL

© guru99.com

CustomerID	CustomerName	Status
1	Google	Active
2	Amazon	Active
3	Apple	Inactive

**Tuple OR Row**  
Total # of rows is **Cardinality**

**Column OR Attributes**  
Total # of column is **Degree**



## Relational Integrity Constraints and Keys

In order to be considered valid, each relation must respect some conditions called *Relational Integrity Constraints*:

- **Key (or Entity) Constraints** - a tuple has to be identified uniquely at least by a *minimal subset of attributes* (this subset is called **key**). If there are more *minimal subsets of attributes*, they are called **candidate keys**. A candidate key that is used to identify a record uniquely in a table is called **primary key**. Key Constraints impose that:
  - A key attribute (or attributes) can't have NULL values
  - In a relation, two tuples cannot have the same values for key attributes
- **Domain Constraints** - attributes can have just a specific set of values
- **Referential Integrity Constraints** - a **foreign key** is a set of attributes that refers to the primary key of another table. The *referential integrity rule* imposes that each foreign key value can be just in 1 of the 2 following states:
  - The foreign key value refers to a primary key value of another table in the RDB
  - Null value (if allowed by the data owner)



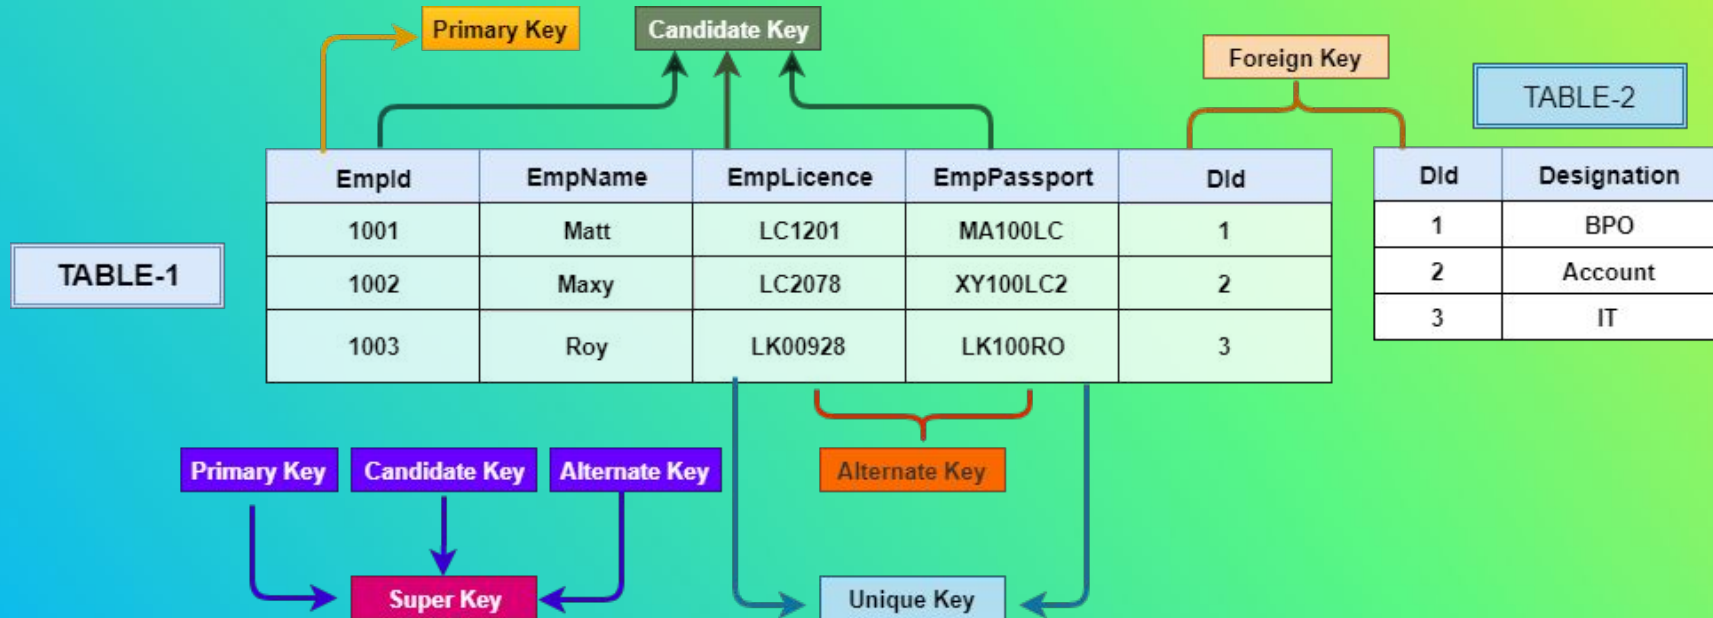


## Types of keys

In the previous slide we mentioned different concepts about keys. Here is a list of all possible keys:

- **Super Key:** a group of single or multiple keys that **identifies rows** in a table
- **Candidate Key:** a set of one or more columns that can **identify a record uniquely** in a table
- **Primary Key:** a candidate key that has been **chosen to identify a record uniquely** in a table
- **Unique Key:** similar to the Primary Key, but it allows one null value in the column
- **Alternate Key:** a possible alternative to the Primary Key
- **Foreign Key:** a reference from Table A to the primary key of another Table B.
  - **The foreign key creates a relationship between 2 or more tables**

# Types of keys - an example





## ACID properties for relational model data

A *database transaction* is a group of tasks that are executed by a **DBMS** (*Database Management System*).

Each database transaction has to observe the ACID properties:

- **Atomicity:** each transaction has to be treaded as an atomic unit. It means that either all the operations (within a transaction) are completed or none. So, it's not possible to have a partially completed database transaction
- **Consistency:** after the transaction execution the database must be remain in a correct state as before
- **Isolation:** each transaction execution is isolated and it doesn't affect any other transaction
- **Durability:** once the transaction is committed, data changes become permanent

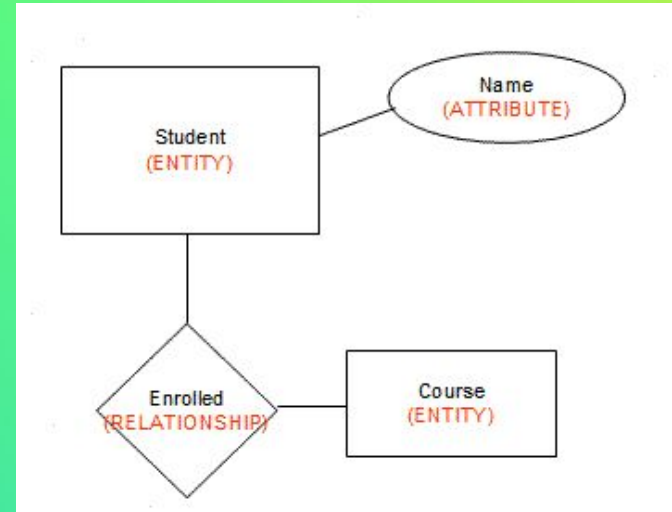


## Entity Relationship Diagram (ERD)

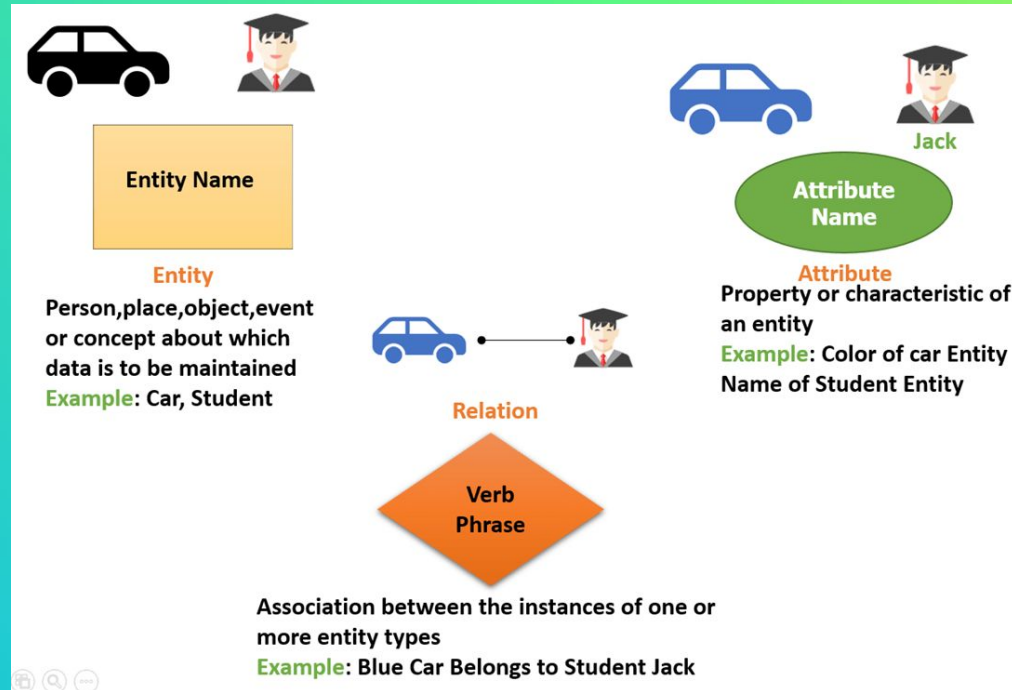
When designing a new database, using some kind of conceptual visualisation tools is really helpful.

One of these tools is the **Entity Relationship Diagram** (ERD) or **ER model**, that is based on 3 basic concepts:

- **Entities** - objects that can have data stored about them → they correspond to the **tables**
- **Attributes** - properties of an Entity
- **Relationship** between entities (as already said, a foreign key creates a relationship between 2 or more tables)



# Entity Relationship Diagram (ERD) - an example



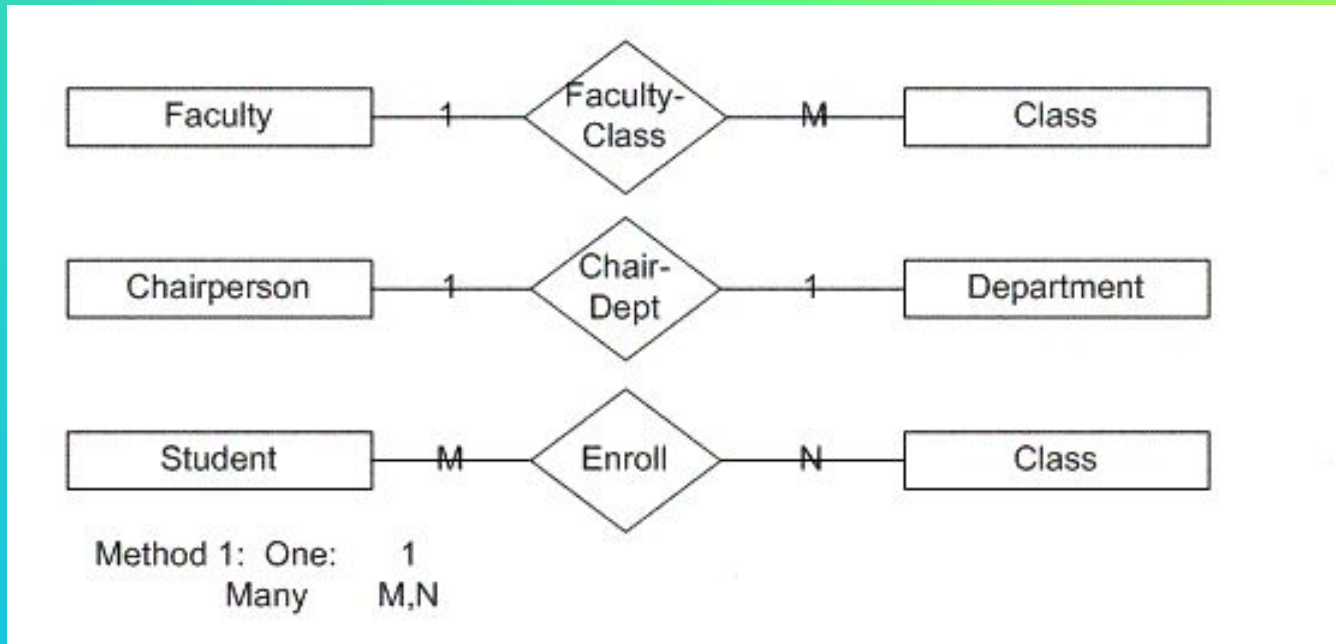


## ER model: types of relationships

ER model's **relationships** can be grouped into 4 main types according to the **relationship cardinality**:

- **One-to-One**
  - One instance of entity A can be associated with at most one instance of entity B
- **One-to-Many**
  - One instance of entity A can be associated with multiple instances of entity B
- **Many-to-One**
  - Multiple instances of entity A can be associated with at most one instance of entity B
- **Many-to-Many**
  - Multiple instances of entity A can be associated with multiple instances of entity B

# ER model: types of relationships - an example





## SQL: database commands

- Basic database creation: `CREATE DATABASE dbnamehere;`
- Create a database if not already existing: `CREATE DATABASE IF NOT EXISTS dbnamehere;`
- Rename a database: `ALTER DATABASE dbnamehere RENAME TO newdbnamehere;`
- Delete a database: `DROP DATABASE dbnamehere;`
- Backup:  
`BACKUP DATABASE dbnamehere;`  
`TO DISK = "filepath";`
- Show available databases: `SHOW DATABASES;`





## SQL: table commands

### Create a table

```
CREATE TABLE table_name (  
    column1 datatype,  
    column2 datatype,  
    ...  
);
```

### Delete a table

```
DROP TABLE table_name;
```

### Update data in the table

```
UPDATE table_name  
SET column_name = value  
WHERE [condition_here];
```

### Add a column to table

```
ALTER TABLE table_name  
ADD new_column datatype;
```

### Remove a column from table

```
ALTER TABLE table_name  
DROP COLUMN column_name;
```



## SQL: SELECT

### Basic SELECT

```
SELECT column_name  
FROM table_name;
```

### SELECT all columns

```
SELECT * FROM table_name;
```

### Select with no duplicates

```
SELECT DISTINCT  
column_name  
FROM table_name;
```

### Select with filter

```
SELECT column_name  
FROM table_name  
WHERE column_name operator value;
```

### Select with a numeric limit

```
SELECT column_name  
FROM table_name  
LIMIT number;
```



## SQL: Logical Operators

**ALL** - True if all of the subquery values meet the condition

```
SELECT ProductName  
FROM Products  
WHERE ProductID = ALL (SELECT ProductID FROM Orders WHERE Quantity = 10);
```

**AND** - TRUE if all the conditions separated by AND are TRUE

```
SELECT * FROM Customers  
WHERE City = "Rome" AND Country = "IT";
```



## SQL: Logical Operators

**ANY** - TRUE if any (at least one) of the subquery values meet the condition

```
SELECT * FROM Products  
WHERE Price > ANY (SELECT Price FROM Products WHERE Price > 50);
```

**BETWEEN** - TRUE if the operand is within the range of comparisons

```
SELECT * FROM Products  
WHERE Price BETWEEN 40 AND 80;
```



## SQL: Logical Operators

**EXISTS** - TRUE if the subquery returns one or more records

```
SELECT SupplierName  
FROM Suppliers  
WHERE EXISTS (SELECT ProductName FROM Products WHERE Products.suppId =  
Suppliers.id AND Price < 50);
```

**IN** - TRUE if the operand is equal to one of a list of expressions

```
SELECT * FROM Customers  
WHERE City IN ('Rome', 'Berlin');
```



## SQL: Logical Operators

**LIKE** - TRUE if the operand matches a pattern

```
SELECT * FROM Customers  
WHERE City LIKE '%in%';
```

**NOT** - Displays a record if the condition(s) is NOT TRUE

```
SELECT * FROM Customers  
WHERE City NOT LIKE '%in%';
```



## SQL: Logical Operators

**OR** - TRUE if any of the conditions separated by OR is TRUE

```
SELECT * FROM Customers  
WHERE City = "Rome" OR Country = "Italy";
```

**SOME** - TRUE if any of the subquery values meet the condition

```
SELECT * FROM Products  
WHERE Price > SOME (SELECT Price FROM Products WHERE Price > 70);
```



## SQL: aggregate functions

**MAX()/MIN()** - Returns the largest/smallest value in a column

```
SELECT MAX|MIN(column_name)
FROM table_name;
```

**SUM()** - Returns the total sum of a numeric column

```
SELECT SUM(age)
FROM students;
```

**AVG()** - The average value for a numeric column

```
SELECT AVG(column_name)
FROM table_name;
```

**COUNT()** - Counts the number of rows where the column is not NULL

```
SELECT COUNT(column_name)
FROM table_name;
```





## SQL: other commands

**AS** - Rename a column or table using an alias

```
SELECT column_name AS 'Alias'  
FROM table_name;
```

**AVG()** - The average value for a numeric column

```
SELECT AVG(column_name)  
FROM table_name;
```

**COUNT()** - Counts the number of rows where the column is not NULL

```
SELECT COUNT(column_name)  
FROM table_name;
```



## SQL: other commands

**DELETE** - Delete rows from a table

```
DELETE FROM table_name  
WHERE column_name = value;
```

**GROUP BY** - Aggregate functions

```
SELECT surname, AVG(age)  
FROM customers  
GROUP BY surname;
```

**HAVING()** - It's like a WHERE for aggregate functions

```
SELECT column_name, COUNT(*)  
FROM table_name  
GROUP BY column_name  
HAVING COUNT(*) > value;
```

**IS NULL** - Test for empty values

```
SELECT column_name  
FROM table_name  
WHERE column_name IS NULL;
```



## SQL: other commands

**ORDER BY** - sort the result set by a particular column either alphabetically or numerically

```
SELECT column_name  
FROM table_name  
ORDER BY column_name ASC | DESC;
```

**TRUNCATE** - removes all data entries from a table in a database, but keeps the table and structure in place


```
TRUNCATE TABLE students;
```

**INSERT** - Add a new row to the table

```
INSERT INTO table_name (column_1, column_2, column_3)  
VALUES (value_1, 'value_2', value_3);
```



# SQL: JOIN - a visual overview





# SQL: INNER JOIN (a kind of AND)

**Records that have matching values in both tables**

```
SELECT name  
FROM customers  
INNER JOIN orders  
ON customers.customer_id = orders.customer_id;
```



# SQL: LEFT JOIN

**Records from the left table that match records in the right table**

```
SELECT name  
FROM customers  
LEFT JOIN orders  
ON customers.customer_id = orders.customer_id;
```



# SQL: RIGHT JOIN

**Records from the right table that match records in the left table**

```
SELECT name  
FROM customers  
RIGHT JOIN orders  
ON customers.customer_id = orders.customer_id;
```



# SQL: FULL JOIN (a kind of OR)

**Records that have a match in the left or right table**

```
SELECT name  
FROM customers  
FULL OUTER JOIN orders  
ON customers.customer_id = orders.customer_id;
```




# SQL: UNION

**UNION** - Combines the result-set (**distinct values**)  
Of two or more SELECT statements

```
SELECT Postcode FROM Customers  
UNION  
SELECT Postcode FROM Suppliers  
ORDER BY Postcode;
```

**UNION ALL** - Combines the result-set (**duplicates values**)  
of two or more SELECT statements

```
SELECT Postcode FROM Customers  
UNION ALL  
SELECT Postcode FROM Suppliers  
ORDER BY Postcode;
```





**DeVeLHOPE**

#codeforimpact

# Introduction to SQL